
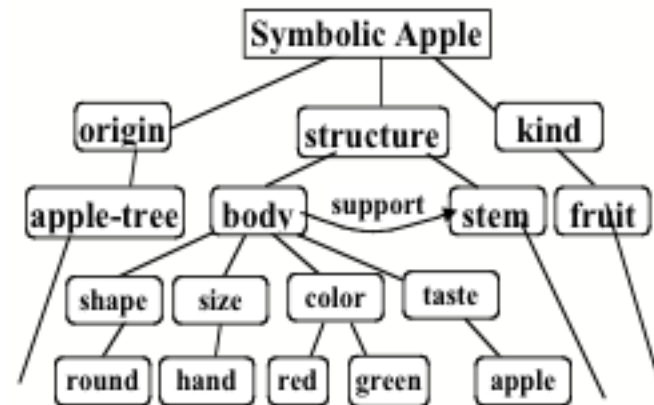


MC4

History of Artificial Intelligence, Neural Networks and Applications

Artificial Intelligence: A Brief History


- Artificial intelligence: *An artificial system developed in computer software, physical hardware, or other context that solves tasks requiring human-like perception, cognition, planning, learning, communication, or physical action* (NDAA, 2019)
- AI was split into two approaches: symbolic and connectionist
 - Symbolic = computation on human-readable representations by logical statements
 - Connectionist = computation on numerical representations by networked connections



Steve Daniels, Medium

Artificial Intelligence: A Brief History

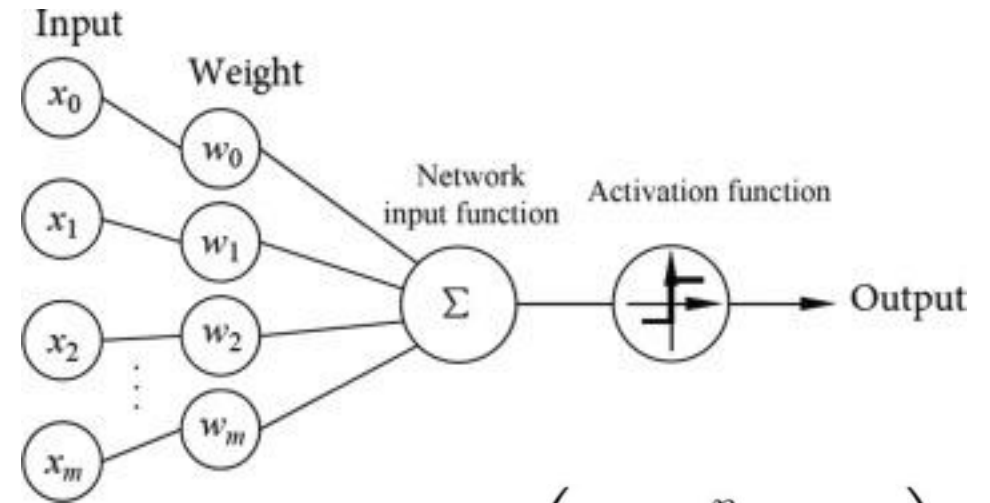
- Artificial intelligence: *An artificial system developed in computer software, physical hardware, or other context that solves tasks requiring human-like perception, cognition, planning, learning, communication, or physical action* (NDAA, 2019)
- AI was split into two approaches: symbolic and connectionist
 - Symbolic = computation on human-readable representations by logical statements
 - Connectionist = computation on numerical representations by networked connections
- In the mid-to-late 20th century, symbolic AI dominated the field
- AI Winters: periods of stalling due to computational or otherwise (70s-80s, 90s)
- In the 2010s, larger connectionist models became more feasible as computing power increased exponentially
- More recently, rise of transformers, some neurosymbolic models




via ExtremeTech

Connectionist Models

- Connectionism uses abstractions of the brain's circuitry known as **Artificial Neural Networks (ANNs)**
- **Perceptron**: early connectionist model, simple architecture
 - **Input layer** receives input values
 - **Weights** applied to each neuron's input
 - **Activation function** allows network to learn nonlinear relationships
 - **Output layer** returns weighted sum of neurons + activation function
- Later, **multilayer perceptron** was introduced to solve more complex problems ([see more on this here](#))
 - The layers in between input and output are known as the **hidden layers**
 - Typically feedforward

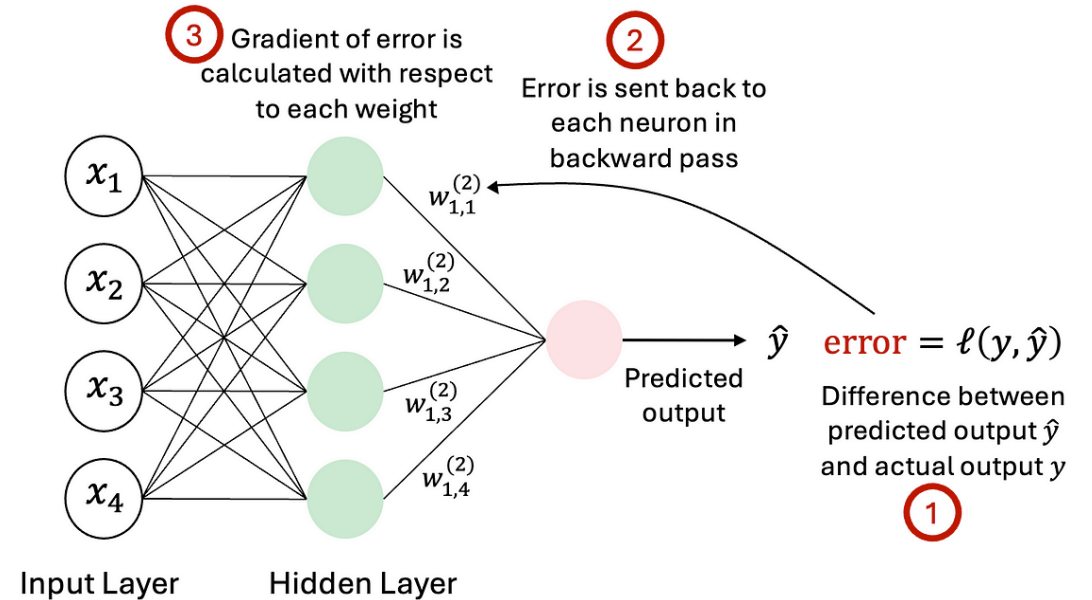


$$f \left(b + \sum_{i=1}^n x_i w_i \right)$$



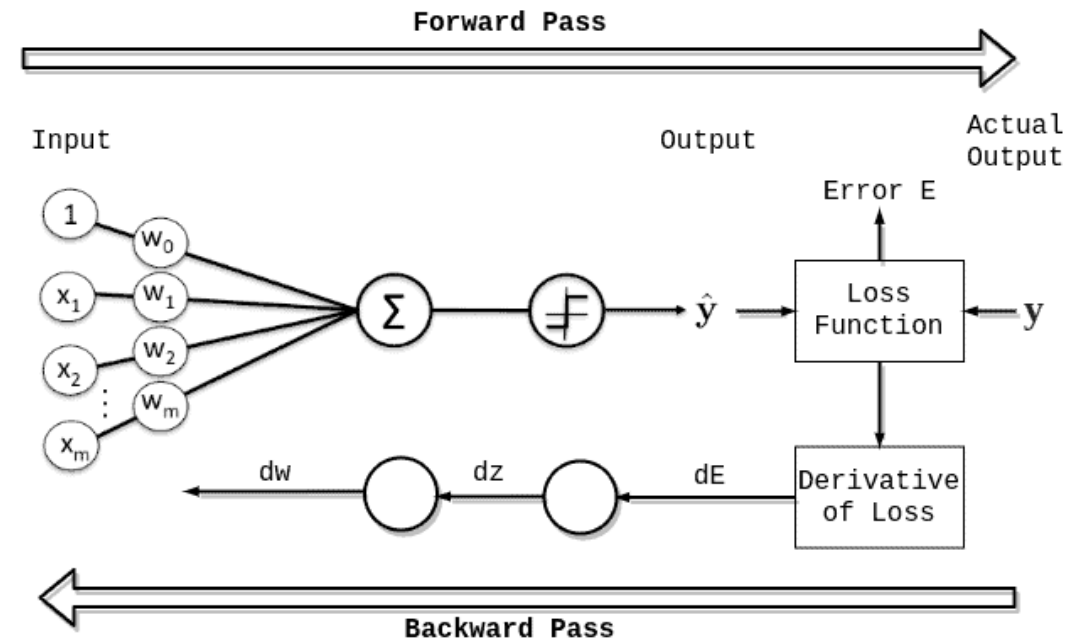


Learning in ANNs vs. Classical ML Models

- Consider the regression problem of fitting weight matrix \mathbf{m} for data approximating $\mathbf{y} = \mathbf{mx} + \mathbf{b}$
- Linear regression typically solves this via approximating \mathbf{m} using least squares
- ANNs learn via:
 - **Propagation:** inputs are passed through the network to obtain an output, and then
 - **Loss Function:** Difference between the output and expected value (error/loss) is calculated
 - **Backpropagation:** Loss used to update the weights of previous layers
- Instead of an exact equation to fit, the model updates its internal weights over the course of several exposures to data



Essential Neural Network Model Training Pipeline

- ANN training typically follows the scheme:
 - Split dataset into **batches**
 - Train neural network on each batch, updating weights on each
 - Do a full pass of all batches N times (each pass is an **epoch**)
- For training evaluation, it is customary to plot the average loss for each epoch



Layers and Operations in ANNs

- **Fully connected layer**
 - All neurons connected
 - Function: full propagation of information between layers
- **Sparsely connected / dropout layer**
 - Few connections between neurons
 - Function: eliminating redundant information, preventing overfitting
- **Convolution layer**
 - Combines local mappings of data via windowed filtering
 - Function: can detect specific features in data (e.g., edges, textures, patterns)
- **Pooling layer**
 - Each $k \times k$ window in data is given a single value
 - Function: Downsamples/aggregates information between layers
- **Flatten layer**
 - Put data into a single 1-dimensional vector
 - Function: Reshape data, combine features
- **Softmax layer**
 - A nonlinear function typically used for mapping output to probabilities between 0-1
 - Typically used in the final layer

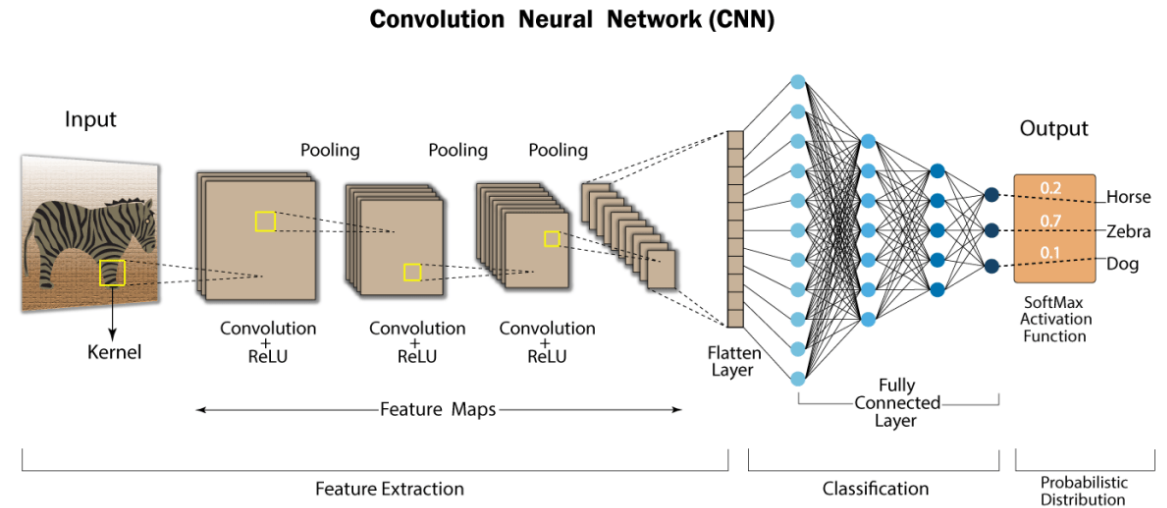



Diagram of a CNN, often used for image processing

Current Hot Topics

- Neural-Representational Alignment
 - How can we build models that align with the human brain?
- Behavioral Alignment
 - How can we build models that perform as well as humans in behavioral tasks (control, perception, memory, etc.)?
- Ethical AI
 - How can we build AI that are aligned with human morals?
- Efficiency
 - How can we build AI that uses fewer computational resources?
- Model Interpretability
 - How can we better understand the hidden layers (internal representations of ANNs)?

Jupyter Notebook Time!

- Today, we will return to our regression problem, but solve with an ANN!
- If you're returning, run **git pull** to update your local repository
- If you're new, clone the git repository for the course:
https://github.com/orbitalhybridization/STARS_ML_MiniCourses